DeepLearning.AI

How to Build Your Career in AI

Collected Insights from Andrew Ng

Founder, DeepLearning.AI

"AI is the new electricity. It will transform and improve all areas of human life."

Andrew Ng

Table of Contents

DeepLearning.AI

Coding AI Is the New Literacy

Today we take it for granted that many people know how to read and write. Someday, I hope, it will be just as common that people know how to write code, specifically for AI.

Several hundred years ago, society didn't view language literacy as a necessary skill. A small number of people learned to read and write, and everyone else let them do the reading and writing. It took centuries for literacy to spread, and now society is far richer for it.

Words enable deep human-to-human communication. Code is the deepest form of human-to-machine communication. As machines become more central to daily life, that communication becomes ever more important.

Traditional software engineering — writing programs that explicitly tell a computer sequences of steps to execute — has been the main path to code literacy. Many introductory programming classes use creating a video game or building a website as examples. But AI, machine learning, and data science offer a new paradigm in which computers extract knowledge from data. This technology offers an even better pathway to coding.

Many Sundays, I buy a slice of pizza from my neighborhood pizza parlor. The gentleman behind the counter has little reason to learn how to build a video game or write his own website software (beyond personal growth and the pleasure of gaining a new skill).

But AI and data science have great value even for a pizza maker. A linear regression model might enable him to better estimate demand so he can optimize the restaurant's staffing and supply chain. He could better predict sales of Hawaiian pizza — my favorite! — so he could make more Hawaiian pies in advance and reduce the amount of time customers had to wait for them.

Uses of AI and data science can be found in almost any situation that produces data. Thus, a wide variety of professions will find more uses for custom AI applications and data-derived insights than for traditional software engineering. This makes literacy in AI-oriented coding even more valuable than traditional coding. It could enable countless individuals to harness data to make their lives richer.

I hope the promise of building basic AI applications, even more than that of building basic traditional software, encourages more people to learn how to code. If society embraces this new form of literacy as it has the ability to read and write, we will all benefit.

Three Steps to Career Growth

The rapid rise of AI has led to a rapid rise in AI jobs, and many people are building exciting careers in this field. A career is a decades-long journey, and the path is not straightforward. Over many years, I've been privileged to see thousands of students, as well as engineers in companies large and small, navigate careers in AI.

Here's a framework for charting your own course.

Three key steps of career growth are **learning foundational skills, working on projects** (to deepen your skills, build a portfolio, and create impact), and **finding a job.** These steps stack on top of each other:

LEARNING   PROJECTS   JOB

Initially, you focus on **learning foundational skills.** Chapters with the 🤖 cover topics about learning foundational technical skills.

After having gained foundational technical skills, you will begin **working on projects.** During this period, you'll also keep learning. Chapters with the 🤖 focus on projects.

Later, you will work on **finding a job.** Throughout this process, you'll continue to learn and work on meaningful projects. Chapters with the 🤖 focus on a job search.

These phases apply in a wide range of professions, but AI involves unique elements. For example:

LEARNING

 Learning foundational skills is a career-long process:

AI is nascent, and many technologies are still evolving. While the foundations of machine learning and deep learning are maturing — and coursework is an efficient way to master them — beyond these foundations, keeping up-to-date with changing technology is more important in AI than fields that are more mature.

PROJECTS

hard to manage because of high iteration

Working on projects often means collaborating with stakeholders who lack expertise in AI:
This can make it challenging to find a suitable project, estimate the project's timeline and return on investment, and set expectations. In addition, the highly iterative nature of AI projects leads to special challenges in project management: How can you come up with a plan for building a system when you don't know in advance how long it will take to achieve the target accuracy? Even after the system has hit the target, further iteration may be necessary to address post-deployment drift.

JOB

Inconsistent opinions on AI skills and jobs roles:
While searching for a job in AI can be similar to searching for a job in other sectors, there are also important differences. Many companies are still trying to figure out which AI skills they need, and how to hire people who have them. Things you've worked on may be significantly different than anything your interviewer has seen, and you're more likely to have to educate potential employers about some elements of your work.

As you go through each step, you should also build a supportive community. Having friends and allies who can help you — and who you strive to help — makes the path easier. This is true whether you're taking your first steps or you've been on the journey for years.

DeepLearning.AI

Learning Technical Skills for a Promising AI Career

LEARNING

In the previous chapter, I introduced three key steps for building a career in AI: learning foundational technical skills, working on projects, and finding a job, all of which is supported by being part of a community. In this chapter, I'd like to dive more deeply into the first step: learning foundational skills.

More research papers have been published on AI than anyone can read in a lifetime. So, when learning, it's critical to prioritize topic selection. I believe the most important topics for a technical career in machine learning are:

Foundational machine learning skills: For example, it's important to understand models such as linear regression, logistic regression, neural networks, decision trees, clustering, and anomaly detection. Beyond specific models, it's even more important to understand the core concepts behind how and why machine learning works, such as bias/variance, cost functions, regularization, optimization algorithms, and error analysis.

Deep learning: This has become such a large fraction of machine learning that it's hard to excel in the field without some understanding of it! It's valuable to know the basics of neural networks, practical skills for making them work (such as hyperparameter tuning), convolutional networks, sequence models, and transformers.

Math relevant to machine learning: Key areas include linear algebra (vectors, matrices, and various manipulations of them) as well as probability and statistics (including discrete and continuous probability, standard probability distributions, basic rules such as independence and Bayes' rule, and hypothesis testing). In addition, exploratory data analysis (EDA) — using visualizations and other methods to systematically explore a dataset — is an underrated skill. I've found EDA particularly useful in data-centric AI development, where analyzing errors and gaining insights can really help drive progress! Finally, a basic intuitive understanding of calculus will also help. The math needed to do machine learning well has been changing. For instance, although some tasks require calculus, improved automatic differentiation software makes it possible to invent and implement new neural network architectures without doing any calculus. This was almost impossible a decade ago.

Software development: While you can get a job and make huge contributions with only machine learning modeling skills, your job opportunities will increase if you can also write good software to implement complex AI systems. These skills include programming fundamentals, data structures (especially those that relate to machine learning, such as data frames), algorithms (including those related to databases and data manipulation), software design, familiarity with Python, and familiarity with key libraries such as TensorFlow or PyTorch, and scikit-learn.

DeepLearning.AI

This is a lot to learn!

Even after you master everything on this list, I hope you'll keep learning and continue to deepen your technical knowledge. I've known many machine learning engineers who benefitted from deeper skills in an application area such as natural language processing or computer vision, or in a technology area such as probabilistic graphical models or building scalable software systems.

How do you gain these skills? There's a lot of good content on the internet, and in theory, reading dozens of web pages could work. But when the goal is deep understanding, reading disjointed web pages is inefficient because they tend to repeat each other, use inconsistent terminology (which slows you down), vary in quality, and leave gaps. That's why a good course — in which a body of material has been organized into a coherent and logical form — is often the most time-efficient way to master a meaningful body of knowledge. When you've absorbed the knowledge available in courses, you can switch over to research papers and other resources.

learn organized course other than disjointed web pages

Finally, no one can cram everything they need to know over a weekend or even a month. Everyone I know who's great at machine learning is a lifelong learner. Given how quickly our field is changing, there's little choice but to keep learning if you want to keep up.

How can you maintain a steady pace of learning for years? If you can cultivate the habit of learning a little bit every week, you can make significant progress with what feels like less effort.

The Best Way to Build a New Habit

One of my favorite books is BJ Fogg's, <u>Tiny Habits: The Small Changes That Change Everything.</u> Fogg explains that the best way to build a new habit is to start small and succeed, rather than start too big and fail. For example, rather than trying to exercise for 30 minutes a day, he recommends aspiring to do just one push-up, and doing it consistently.

This approach may be helpful to those of you who want to spend more time studying. If you start by holding yourself accountable for watching, say, 10 seconds of an educational video every day — and you do so consistently — the habit of studying daily will grow naturally. Even if you learn nothing in that 10 seconds, you're establishing the habit of studying a little every day. On some days, maybe you'll end up studying for an hour or longer.

DeepLearning.AI

Should You Learn Math to Get a Job in AI?

How much math do you need to know to be a machine learning engineer?

Is math a foundational skill for AI? It's always nice to know more math! But there's so much to learn that, realistically, it's necessary to prioritize. Here's how you might go about strengthening your math background.

To figure out what's important to know, I find it useful to ask what you need to know to make the decisions required for the work you want to do. At DeepLearning.AI, we frequently ask, "What does someone need to know to accomplish their goals?" The goal might be building a machine learning model, architecting a system, or passing a job interview.

ai more mature, less understanding of math enable us to work

Understanding the math behind algorithms you use is often helpful, since it enables you to debug them. But the depth of knowledge that's useful changes over time. As machine learning techniques mature and become more reliable and turnkey, they require less debugging, and a shallower understanding of the math involved may be sufficient to make them work.

For instance, in an earlier era of machine learning, linear algebra libraries for solving linear systems of equations (for linear regression) were immature. I had to understand how these libraries worked so I could choose among different libraries and avoid numerical roundoff pitfalls. But this became less important as numerical linear algebra libraries matured.

understanding math behind it will help us understand and make better decisions.

Deep learning is still an emerging technology, so when you train a neural network and the optimization algorithm struggles to converge, understanding the math behind gradient descent, momentum, and the Adam optimization algorithm will help you make better decisions. Similarly, if your neural network does something funny — say, it makes bad predictions on images of a certain resolution, but not others — understanding the math behind neural network architectures puts you in a better position to figure out what to do.

learn by intrinsic curiosity

Of course, I also encourage learning driven by curiosity. If something interests you, go ahead and learn it regardless of how useful it might turn out to be! Maybe this will lead to a creative spark or technical breakthrough.

Scoping Successful AI Projects

PROJECTS

One of the most important skills of an AI architect is the ability to identify ideas that are worth working on. These next few chapters will discuss finding and working on projects so you can gain experience and build your portfolio.

Over the years, I've had fun applying machine learning to manufacturing, healthcare, climate change, agriculture, ecommerce, advertising, and other industries. How can someone who's not an expert in all these sectors find meaningful projects within them? Here are five steps to help you scope projects.

Step 1

Identify a business problem (not an AI problem). I like to find a domain expert and ask, "What are the top three things that you wish worked better? Why aren't they working yet?" For example, if you want to apply AI to climate change, you might discover that power-grid operators can't accurately predict how much power intermittent sources like wind and solar might generate in the future.

Step 2

Brainstorm AI solutions. When I was younger, I used to execute on the first idea I was excited about. Sometimes this worked out okay, but sometimes I ended up missing an even better idea that might not have taken any more effort to build. Once you understand a problem, you can brainstorm potential solutions more efficiently. For instance, to predict power generation from intermittent sources, we might consider using satellite imagery to map the locations of wind turbines more accurately, using satellite imagery to estimate the height and generation capacity of wind turbines, or using weather data to better predict cloud cover and thus solar irradiance. Sometimes there isn't a good AI solution, and that's okay too.

DeepLearning.AI

Step 3

Assess the feasibility and value of potential solutions. You can determine whether an approach is technically feasible by looking at published work, what competitors have done, or perhaps building a quick proof of concept implementation. You can determine its value by consulting with domain experts (say, power-grid operators, who can advise on the utility of the potential solutions mentioned above).

Step 4

Determine milestones. Once you've deemed a project sufficiently valuable, the next step is to determine the metrics to aim for. This includes both machine learning metrics (such as accuracy) and business metrics (such as revenue). Machine learning teams are often most comfortable with metrics that a learning algorithm can optimize. But we may need to stretch outside our comfort zone to come up with business metrics, such as those related to user engagement, revenue, and so on. Unfortunately, not every business problem can be reduced to optimizing test set accuracy! If you aren't able to determine reasonable milestones, it may be a sign that you need to learn more about the problem. A quick proof of concept can help supply the missing perspective.

Step 5

Budget for resources. Think through everything you'll need to get the project done including data, personnel, time, and any integrations or support you may need from other teams. For example, if you need funds to purchase satellite imagery, make sure that's in the budget.

Working on projects is an iterative process. If, at any step, you find that the current direction is infeasible, return to an earlier step and proceed with your new understanding. Is there a domain that excites you where AI might make a difference? I hope these steps will guide you in exploring it through project work — even if you don't yet have deep expertise in that field. AI won't solve every problem, but as a community, let's look for ways to make a positive impact wherever we can.

DeepLearning.AI

Finding Projects that Complement Your Career Goals

PROJECTS

It goes without saying that we should only work on projects that are responsible, ethical, and beneficial to people. But those limits leave a large variety to choose from. In the previous chapter, I wrote about how to identify and scope AI projects. This chapter and the next have a slightly different emphasis: picking and executing projects with an eye toward career development.

A fruitful career will include many projects, hopefully growing in scope, complexity, and impact over time. **Thus, it is fine to start small.** Use early projects to learn and gradually step up to bigger projects as your skills grow.

When you're starting out, don't expect others to hand great ideas or resources to you on a platter. Many people start by working on small projects in their spare time. With initial successes — even small ones — under your belt, your growing skills increase your ability to come up with better ideas, and it becomes easier to persuade others to help you step up to bigger projects.

What if you don't have any project ideas? Here are a few ways to generate them:

✓ **Join existing projects.** If you find someone else with an idea, ask to join their project.

✓ **Keep reading and talking to people.** I come up with new ideas whenever I spend a lot of time reading, taking courses, or talking with domain experts. I'm confident that you will, too.

✓ **Focus on an application area.** Many researchers are trying to advance basic AI technology — say, by inventing the next generation of transformers or further scaling up language models — so, while this is an exciting direction, it is also very hard. But the variety of applications to which machine learning has not yet been applied is vast! I'm fortunate to have been able to apply neural networks to everything from autonomous helicopter flight to online advertising, partly because I jumped in when relatively few people were working on those applications. If your company or school cares about a particular application, explore the possibilities for machine learning. That can give you a first look at a potentially creative application — one where you can do unique work — that no one else has done yet.

DeepLearning.AI

✔ **Develop a side hustle.** Even if you have a full-time job, a fun project that may or may not develop into something bigger can stir the creative juices and strengthen bonds with collaborators. When I was a full-time professor, working on online education wasn't part of my "job" (which was doing research and teaching classes). It was a fun hobby that I often worked on out of passion for education. My early experiences in recording videos at home helped me later in working on online education in a more substantive way. Silicon Valley abounds with stories of startups that started as side projects. As long as it doesn't create a conflict with your employer, these projects can be a stepping stone to something significant.

Given a few project ideas, which one should you jump into? Here's a quick checklist of factors to consider:

✔ **Will the project help you grow technically?** Ideally, it should be challenging enough to stretch your skills but not so hard that you have little chance of success. This will put you on a path toward mastering ever-greater technical complexity.

✔ **Do you have good teammates to work with?** If not, are there people you can discuss things with? We learn a lot from the people around us, and good collaborators will have a huge impact on your growth.

✔ **Can it be a stepping stone?** If the project is successful, will its technical complexity and/ or business impact make it a meaningful stepping stone to larger projects? If the project is bigger than those you've worked on before, there's a good chance it could be such a stepping stone.

Finally, avoid analysis paralysis. It doesn't make sense to spend a month deciding whether to work on a project that would take a week to complete. You'll work on multiple projects over the course of your career, so you'll have ample opportunity to refine your thinking on what's worthwhile. Given the huge number of possible AI projects, rather than the conventional "ready, aim, fire" approach, you can accelerate your progress with "ready, fire, aim."

Ready, Fire, Aim

Working on projects requires making tough choices about what to build and how to go about it. Here are two distinct styles:

- ✓ **Ready, Aim, Fire:** Plan carefully and carry out careful validation. Commit and execute only when you have a high degree of confidence in a direction.

- ✓ **Ready, Fire, Aim:** Jump into development and start executing. This allows you to discover problems quickly and pivot along the way if necessary.

Say you've built a customer-service chatbot for retailers, and you think it could help restaurants, too. Should you take time to study the restaurant market before starting development, moving slowly but cutting the risk of wasting time and resources? Or jump in right away, moving quickly and accepting a higher risk of pivoting or failing?

Both approaches have their advocates, and the best choice depends on the situation.

Ready, Aim, Fire tends to be superior when the cost of execution is high and a study can shed light on how useful or valuable a project could be. For example, if you can brainstorm a few other use cases (restaurants, airlines, telcos, and so on) and evaluate these cases to identify the most promising one, it may be worth taking the extra time before committing to a direction.

Ready, Fire, Aim tends to be better if you can execute at low cost and, in doing so, determine whether the direction is feasible and discover tweaks that will make it work. For example, if you can build a prototype quickly to figure out if users want the product, and if canceling or pivoting after a small amount of work is acceptable, then it makes sense to consider jumping in quickly. When taking a shot is inexpensive, it also makes sense to take many shots. In this case, the process is actually Ready, Fire, Aim, Fire, Aim, Fire, Aim, Fire.

After agreeing upon a project direction, when it comes to building a machine learning model that's part of the product, I have a bias toward Ready, Fire, Aim. Building models is an iterative process. For many applications, the cost of training and conducting error analysis is not prohibitive. Furthermore, it is very difficult to carry out a study that will shed light on the appropriate model, data, and hyperparameters. So it makes sense to build an end-to-end system quickly and revise it until it works well.

for model training, ng suggests ready fire aim do and pivot/tweet

But when committing to a direction means making a costly investment or entering a one-way door (meaning a decision that's hard to reverse), it's often worth spending more time in advance to make sure it really is a good idea.

DeepLearning.AI

Building a Portfolio of Projects that Shows Skill Progression

PROJECTS

Over the course of a career, you're likely to work on projects in succession, each growing in scope and complexity. For example:

1. Class projects:

The first few projects might be narrowly scoped homework assignments with predetermined right answers. These are often great learning experiences!

2. Personal projects

You might go on to work on small-scale projects either alone or with friends. For instance, you might re-implement a known algorithm, apply machine learning to a hobby (such as predicting whether your favorite sports team will win), or build a small but useful system at work in your spare time (such as a machine learning-based script that helps a colleague automate some of their work). Participating in competitions such as those organized by Kaggle is also one way to gain experience.

3. Creating value

Eventually, you will gain enough skill to build projects in which others see more tangible value. This opens the door to more resources. For example, rather than developing machine learning systems in your spare time, it might become part of your job, and you might gain access to more equipment, compute time, labeling budget, or head count.

4. Rising scope and complexity

Successes build on each other, opening the door to more technical growth, more resources, and increasingly significant project opportunities.

DeepLearning.AI

Each project is only one step on a longer journey, hopefully one that has a positive impact. In addition:

Don't worry about starting too small. One of my first machine learning research projects involved training a neural network to see how well it could mimic the sin(x) function. It wasn't very useful, but was a great learning experience that enabled me to move on to bigger projects.

Communication is key. You need to be able to explain your thinking if you want others to see the value in your work and trust you with resources that you can invest in larger projects. To get a project started, communicating the value of what you hope to build will help bring colleagues, mentors, and managers onboard — and help them point out flaws in your reasoning. After you've finished, the ability to explain clearly what you accomplished will help convince others to open the door to larger projects.

Leadership isn't just for managers. When you reach the point of working on larger AI projects that require teamwork, your ability to lead projects will become more important, whether or not you are in a formal position of leadership. Many of my friends have successfully pursued a technical rather than managerial career, and their ability to help steer a project by applying deep technical insights — for example, when to invest in a new technical architecture or collect more data of a certain type — allowed them to grow as leaders and also helped significantly improve the project.

> Building a portfolio of projects, especially one that shows progress over time from simple to complex undertakings, will be a big help when it comes to looking for a job.

# A Simple Framework for Starting Your AI Job Search

JOBS

Finding a job has a few predictable steps that include <mark>selecting the companies to which you want to apply, preparing for interviews, and finally picking a role and negotiating a salary and benefits.</mark> In this chapter, I'd like to focus on a framework that's useful for many job seekers in AI, especially those who are entering AI from a different field.

If you're considering your next job, ask yourself:

- ✓ **Are you switching roles?** For example, if you're a software engineer, university student, or physicist who's looking to become a machine learning engineer, that's a role switch.
- ✓ **Are you switching industries?** For example, if you work for a healthcare company, financial services company, or a government agency and want to work for a software company, that's a switch in industries.

A product manager at a tech startup who becomes a data scientist at the same company (or a different one) has switched roles. A marketer at a manufacturing firm who becomes a marketer in a tech company has switched industries. An analyst in a financial services company who becomes a machine learning engineer in a tech company has switched both roles and industries.

If you're looking for your first job in AI, you'll probably find switching either roles or industries easier than doing both at the same time. Let's say you're the analyst working in financial services:

- ✓ If you find a data science or machine learning job in financial services, you can continue to use your domain-specific knowledge while gaining knowledge and expertise in AI. After working in this role for a while, you'll be better positioned to switch to a tech company (if that's still your goal).

- ✓ Alternatively, if you become an analyst in a tech company, you can continue to use your skills as an analyst but apply them to a different industry. Being part of a tech company also makes it much easier to learn from colleagues about practical challenges of AI, key skills to be successful in AI, and so on.

　　　　　　　　　　　　　　　　　　　　　　　　　　　　　　　DeepLearning.AI

If you're considering a role switch, a startup can be an easier place to do it than a big company. While there are exceptions, startups usually don't have enough people to do all the desired work. If you're able to help with AI tasks — even if it's not your official job — your work is likely to be appreciated. This lays the groundwork for a possible role switch without needing to leave the company. In contrast, in a big company, a rigid reward system is more likely to reward you for doing your job well (and your manager for supporting you in doing the job for which you were hired), but it's not as likely to reward contributions outside your job's scope.

After working for a while in your desired role and industry (for example, a machine learning engineer in a tech company), you'll have a good sense of the requirements for that role in that industry at a more senior level. You'll also have a network within that industry to help you along. So future job searches — if you choose to stick with the role and industry — likely will be easier.

When changing jobs, you're taking a step into the unknown, particularly if you're switching either roles or industries. One of the most underused tools for becoming more familiar with a new role and/or industry is the informational interview. I'll share more about that in the next chapter.

I'm grateful to Salwa Nur Muhammad, CEO of FourthBrain (a DeepLearning.AI affiliate), for providing some of the ideas presented in this chapter.

Overcoming Uncertainty

There's a lot we don't know about the future: When will we cure Alzheimer's disease? Who will win the next election? Or, in a business context, how many customers will we have next year?

With so many changes going on in the world, many people are feeling stressed about the future, especially when it comes to finding a job. I have a practice that helps me regain a sense of control. Faced with uncertainty, I try to:

1 Make a list of plausible scenarios, acknowledging that I don't know which will come to pass.

2 Create a plan of action for each scenario.

3 Start executing actions that seem reasonable.

**4** Review scenarios and plans periodically as the future comes into focus.

For example, during the Covid-19 pandemic back in March 2020, I did this scenario planning exercise. I imagined quick (three months), medium (one year), and slow (two years) recoveries from Covid-19 and made plans for managing each case. These plans have helped me prioritize where I can.

The same method can apply to personal life, too. If you're not sure you'll pass an exam, get a job offer, or be granted a visa — all of which can be stressful — you can write out what you'd do in each of the likely scenarios. Thinking through the possibilities and following through on plans can help you navigate the future effectively no matter what it brings.

Bonus: With training in AI and statistics, you can calculate a probability for each scenario. I'm a fan of the Superforecasting methodology, in which the judgments of many experts are synthesized into a probability estimate.

DeepLearning.AI

Using Informational Interviews to Find the Right Job

JOBS

If you're preparing to switch roles (say, taking a job as a machine learning engineer for the first time) or industries (say, working in an AI tech company for the first time), there's a lot about your target job that you probably don't know. A technique known as informational interviewing is a great way to learn.

An informational interview involves finding someone in a company or role you'd like to know more about and informally interviewing them about their work. Such conversations are separate from searching for a job. In fact, it's helpful to interview people who hold positions that align with your interests well before you're ready to kick off a job search.

✓ Informational interviews are particularly relevant to AI. Because the field is evolving, many companies use job titles in inconsistent ways. In one company, data scientists might be expected mainly to analyze business data and present conclusions on a slide deck. In another, they might write and maintain production code. An informational interview can help you sort out what the AI people in a particular company actually do.

✓ With the rapid expansion of opportunities in AI, many people will be taking on an AI job for the first time. In this case, an informational interview can be invaluable for learning what happens and what skills are needed to do the job well. For example, you can learn what algorithms, deployment processes, and software stacks a particular company uses. You may be surprised — if you're not already familiar with the data-centric AI movement — to learn how much time most machine learning engineers spend iteratively cleaning datasets.

Prepare for informational interviews by researching the interviewee and company in advance, so you can arrive with thoughtful questions. You might ask:

✓ What do you do in a typical week or day?
✓ What are the most important tasks in this role?
✓ What skills are most important for success?
✓ How does your team work together to accomplish its goals?
✓ What is the hiring process?
✓ Considering candidates who stood out in the past, what enabled them to shine?

Finding someone to interview isn't always easy, but many people who are in senior positions today received help when they were new from those who had entered the field ahead of them, and many are eager to pay it forward. If you can reach out to someone who's already in your network — perhaps a friend who made the transition ahead of you or someone who attended the same school as you — that's great! Meetups such as Pie & AI can also help you build your network.

Finally, be polite and professional, and thank the people you've interviewed. And when you get a chance, please pay it forward as well and help someone coming up after you. If you receive a request for an informational interview from someone in the DeepLearning.AI community, I hope you'll lean in to help them take a step up! If you're interested in learning more about informational interviews, I recommend this article from the UC Berkeley Career Center.

I've mentioned a few times the importance of your network and community. People you've met, beyond providing valuable information, can also play an invaluable role by referring you to potential employers.

Finding the Right AI Job for You

JOBS

In this chapter, I'd like to discuss some fine points of finding a job.

The typical job search follows a fairly predictable path.

- ✓ Research roles and companies online or by talking to friends.
- ✓ Optionally, arrange informal informational interviews with people in companies that appeal to you.
- ✓ Either apply directly or, if you can, get a referral from someone on the inside.
- ✓ Interview with companies that give you an invitation.
- ✓ Receive one or more offers and pick one. Or, if you don't receive an offer, ask for feedback from the interviewers, human resources staff, online discussion boards, or anyone in your network who can help you plot your next move.

Although the process may be familiar, every job search is different. Here are some tips to increase the odds you'll find a position that supports your thriving career and enables you to keep growing.

Pay attention to the fundamentals. A compelling resume, portfolio of technical projects, and a strong interview performance will unlock doors. Even if you have a referral from someone in a company, a resume and portfolio will be your first contact with many people who don't already know about you. Update your resume and make sure it clearly presents your education and experience relevant to the role you want. Customize your communications with each company to explain why you're a good fit. Before an interview, ask the recruiter what to expect. Take time to review and practice answers to common interview questions, brush up key skills, and study technical materials to make sure they are fresh in your mind. Afterward, take notes to help you remember what was said.

Proceed respectfully and responsibly. Approach interviews and offer negotiations with a win-win mindset. Outrage spreads faster than reasonableness on social media, so a story about how an employer underpaid someone gets amplified, whereas stories about how an employer treated someone fairly do not. The vast majority of employers are ethical and fair, so don't let stories about the small fraction of mistreated individuals sway your approach. If you're leaving a job, exit gracefully. Give your employer ample notice, give your full effort through your last hour on the job, transition unfinished business as best you can, and leave in a way that honors the responsibilities you were entrusted with.

Choose who to work with. It's tempting to take a position because of the projects you'll work on. But the teammates you'll work with are at least equally important. We're influenced by people around us, so your colleagues will make a big difference. For example, if your friends smoke, the odds increase that you, too, will smoke. I don't know of a study that shows this, but I'm pretty sure that if most of your colleagues work hard, learn continuously, and build AI to benefit all people, you're likely to do the same. (By the way, some large companies won't tell you who your teammates will be until you've accepted an offer. In this case, be persistent and keep pushing to identify and speak with potential teammates. Strict policies may make it impossible to accommodate you, but in my mind, that increases the risk of accepting the offer, as it increases the odds you'll end up with a manager or teammates who aren't a good fit.)

Get help from your community. Most of us go job hunting only a small number of times in our careers, so few of us get much practice at doing it well. Collectively, though, people in your immediate community probably have a lot of experience. Don't be shy about calling on them. Friends and associates can provide advice, share inside knowledge, and refer you to others who may help. I got a lot of help from supportive friends and mentors when I applied for my first faculty position, and many of the tips they gave me were very helpful.

I know that the job-search process can be intimidating. Instead of viewing it as a great leap, consider an incremental approach. Start by identifying possible roles and conducting a handful of informational interviews. If these conversations tell you that you have more learning to do before you're ready to apply, that's great! At least you have a clear path forward. The most important part of any journey is to take the first step, and that step can be a small one.

DeepLearning.AI

Keys to Building a Career in AI

The path to career success in AI is more complex than what I can cover in one short eBook.

Hopefully the previous chapters will give you momentum to move forward.

Here are additional things to think about as you plot your path to success:

1. Teamwork:

When we tackle large projects, we succeed better by working in teams than individually. The ability to collaborate with, influence, and be influenced by others is critical. Thus, interpersonal and communication skills really matter. (I used to be a pretty <u>bad communicator,</u> by the way.)

2. Networking:

I hate networking! As an introvert, having to go to a party to smile and shake as many hands as possible is an activity that borders on horrific. I'd much rather stay home and read a book. Nonetheless, I'm fortunate to have found many genuine friends in AI; people I would gladly go to bat for and who I count on as well. No person is an island, and having a strong professional network can help propel you forward in the moments when you need help or advice. In lieu of networking, I've found it more helpful to think about building up a community. So instead of trying to build up my personal network, I focus instead on building up the communities that I'm part of. This has the side effect of helping me meet more people and make friends as well.

DeepLearning.AI

3. Job search

Of all the steps in building a career, this one tends to receive the most attention. Unfortunately, there is a lot of bad advice about this on the internet. (For example, many articles urge taking an adversarial attitude toward potential employers, which I don't think is helpful.) Although it may seem like finding a job is the ultimate goal, it's just one small step in the long journey of a career.

4. Personal discipline

Few people will know whether you spend your weekends learning, or binge watching TV — but they will notice the difference over time. Many successful people develop good habits in eating, exercise, sleep, personal relationships, work, learning, and self-care. Such habits help them move forward while staying healthy.

5. Altruism

I find that people who aim to lift others during every step of their own journey often achieve better outcomes for themselves. How can we help others even as we build an exciting career for ourselves?

CHAPTER 11
Overcoming Imposter Syndrome

Before we dive into the final chapter of this book, I'd like to address the serious matter of newcomers to AI sometimes experiencing imposter syndrome, where someone — regardless of their success in the field — wonders if they're a fraud and really belong in the AI community. I want to make sure this doesn't discourage you or anyone else from growing in AI.

Let me be clear: If you want to be part of the AI community, then I welcome you with open arms. If you want to join us, you fully belong with us!

An estimated 70 percent of people experience some form of imposter syndrome at some point. Many talented people have spoken publicly about this experience, including former Facebook COO Sheryl Sandberg, U.S. first lady Michelle Obama, actor Tom Hanks, and Atlassian co-CEO Mike Cannon-Brookes. It happens in our community even among accomplished people. If you've never experienced this yourself, that's great! I hope you'll join me in encouraging and welcoming everyone who wants to join our community.

AI is technically complex, and it has its fair share of smart and highly capable people. But it is easy to forget that to become good at anything, the first step is to suck at it. If you've succeeded at sucking at AI — congratulations, you're on your way!

I once struggled to understand the math behind linear regression. I was mystified when logistic regression performed strangely on my data, and it took me days to find a bug in my implementation of a basic neural network. Today, I still find many research papers challenging to read, and I recently made an obvious mistake while tuning a neural network hyperparameter (that fortunately a fellow engineer caught and fixed).

So if you, too, find parts of AI challenging, it's okay. We've all been there. I guarantee that everyone who has published a seminal AI paper struggled with similar technical challenges at some point.

Here are some things that can help.

✓ **Do you have supportive mentors or peers?** If you don't yet, attend Pie & AI or other events, use discussion boards, and work on finding some. If your mentors or manager don't support your growth, find ones who do. I'm also working on how to grow a supportive AI community and hope to make finding and giving support easier for everyone.

✓ **No one is an expert at everything.** Recognize what you do well. If what you do well is understand and explain to your friends one-tenth of the articles in The Batch, then you're on your way! Let's work on getting you to understand two-tenths of the articles.

My three-year-old daughter (who can barely count to 12) regularly tries to teach things to my one-year-old son. No matter how far along you are — if you're at least as knowledgeable as a three-year-old — you can encourage and lift up others behind you. Doing so will help you, too, as others behind you will recognize your expertise and also encourage you to keep developing. When you invite others to join the AI community, which I hope you will do, it also reduces any doubts that you are already one of us.

AI is such an important part of our world that I would like everyone who wants to be part of it to feel at home as a member of our community. Let's work together to make it happen.

DeepLearning.AI

Final Thoughts

Make Every Day Count

Every year on my birthday, I get to thinking about the days behind and those that may lie ahead.

Maybe you're good at math; I'm sure you'll be able to answer the following question via a quick calculation. But let me ask you a question, and please *answer from your gut, without calculating.*

How many days is a typical human lifespan?

20,000 days	100,000 days
1 million days	5 million days

When I ask friends, many choose a number in the hundreds of thousands. (Many others can't resist calculating the answer, to my annoyance!)

When I was a grad student, I remember plugging my statistics into a mortality calculator to figure out my life expectancy. The calculator said I could expect to live a total of 27,649 days. It struck me how small this number is. I printed it in a large font and pasted it on my office wall as a daily reminder.

That's all the days we have to spend with loved ones, learn, build for the future, and help others. Whatever you're doing today, is it worth 1/30,000 of your life?

DeepLearning.AI